

Principles of Data Science

Name: Revati Jaidatta Chavare (22533031)

Section-1 Data Understanding and exploration

Importing required libraries..

In [390]:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%config InlineBackend.figure_format='retina'
sns.set(
    rc={ "figure.figsize": (14,6) },
    style="ticks", context="notebook", font_scale=1.2
)
```

Loading dataset and storing it into 'advert'

In [391]:

```
advert=pd.read_csv('adverts.csv')
```

In [392]:

```
advert.head(5)
```

Out[392]:

	public_reference	mileage	reg_code	standard_colour	standard_make	standard_model	vehicle_condition	year_of_registration	price	body_type	crossover_car_and_van
0	202006039777689	0.0	NaN	Grey	Volvo	XC90	NEW	NaN	73970	SUV	
1	202007020778260	108230.0	61	Blue	Jaguar	XF	USED	2011.0	7000	Saloon	
2	202007020778474	7800.0	17	Grey	SKODA	Yeti	USED	2017.0	14000	SUV	
3	202007080986776	45000.0	16	Brown	Vauxhall	Mokka	USED	2016.0	7995	Hatchback	
4	202007161321269	64000.0	64	Grey	Land Rover	Range Rover Sport	USED	2015.0	26995	SUV	

In [393]:

```
advert.columns
```

Out[393]:

```
Index(['public_reference', 'mileage', 'reg_code', 'standard_colour',
      'standard_make', 'standard_model', 'vehicle_condition',
      'year_of_registration', 'price', 'body_type', 'crossover_car_and_van',
      'fuel_type'],
      dtype='object')
```

In [394]:

```
advert.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 402005 entries, 0 to 402004
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   public_reference                      402005 non-null int64
1   mileage                              401878 non-null float64
2   reg_code                             370148 non-null object
3   standard_colour                       396627 non-null object
4   standard_make                         402005 non-null object
5   standard_model                       402005 non-null object
6   vehicle_condition                    402005 non-null object
7   year_of_registration                 368694 non-null float64
8   price                                402005 non-null int64
9   body_type                            401168 non-null object
10  crossover_car_and_van                 402005 non-null bool
11  fuel_type                             401404 non-null object
dtypes: bool(1), float64(2), int64(2), object(7)
memory usage: 34.1+ MB
```

1.1 Identifying missing/ null values

Checking for null values in the raw dataset, using `isna()` function and `sum()` function to get the total number of null values in each column.

In [395]:

```
advert.isna().sum()
```

Out[395]:

```
public_reference      0
mileage               127
reg_code             31857
standard_colour       5378
standard_make         0
standard_model        0
vehicle_condition     0
year_of_registration  33311
price                 0
body_type             837
crossover_car_and_van 0
fuel_type             601
dtype: int64
```

Here, we can see `mileage`, `reg_code`, `standard_colour`, `year_of_registration`, `body_type` and `fuel_type` contains null values.


Using heatmap to visualise null values in the dataset.

In [396]:

```
sns.heatmap(advert.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[396]:

<AxesSubplot:>



1.2 Identifying outliers or noise


An outlier is a distant observation in the dataset, it basically falls outside of the usual observation and can be easily spotted using boxplot or scatterplot. Here, we are looking for outliers in price, mileage and year of registration using boxplot and scatterplot respectively.

In [397]:

```
sns.boxplot(data=advert['price'])
```

Out[397]:

<AxesSubplot:>




In [398]:

```
sns.scatterplot(data=advert['year_of_registration'])
```

Out[398]:

<AxesSubplot:ylabel='year_of_registration'>




In [399]:

```
sns.boxplot(data=advert['mileage'])
```

Out[399]:

<AxesSubplot:>



After carefully observing the plots we are now sure that there are outliers in mileage, price and year of registration and we'll have to deal with it in further steps.

1.3 Identifying features

Features considered:

Numerical Features:

1. Mileage
2. Year of registration

Categorical Features:

1. Vehicle condition
2. Body type
3. Fuel type

Many features in the dataset are not that important and we can drop those features, and work only on potential features to predict the price of a car.

In [400]:

```
advert.columns
```

Out[400]:

```
Index(['public_reference', 'mileage', 'reg_code', 'standard_colour',
      'standard_make', 'standard_model', 'vehicle_condition',
      'year_of_registration', 'price', 'body_type', 'crossover_car_and_van',
      'fuel_type'],
      dtype='object')
```

In [401]:

```
advert.corr()['price']
```

Out[401]:

```
public_reference    -0.052344
mileage             -0.160204
year_of_registration  0.102341
price               1.000000
crossover_car_and_van  0.010402
Name: price, dtype: float64
```

A quick look at the correlation between price and other numerical features in the dataset. This clearly shows that year-of_registration could play a vital role in determining the price of a car, it has a significantly high correlation with price comparative to other features.

In [459]:

```
Total_no_of_cars=advert.groupby(['standard_make', 'standard_model']).size().reset_index().rename(columns={0: 'Total number of cars'})
```

In [460]:

Total_no_of_cars

Out[460]:

	standard_make	standard_model	Total number of cars
0	AC	Cobra	3
1	AK	Cobra	1
2	Abarth	124 Spider	61
3	Abarth	500	109
4	Abarth	500C	27
...
1212	Westfield	Se	1
1213	Westfield	Sei	1
1214	Westfield	Sport	1
1215	Wolseley	6/110	1

Using `group_by()` function, here we got the total number of cars of each `standard_make` and `standard_model` in the dataset.

Section-2 Data processing

2.1 Handling missing/ null values

As mentioned earlier, we can drop some features and work on potential features. So, here we are dropping `public_reference`, `reg_code`, `standard_colour`, `standard_make`, `standard_model` and `crossover_car_and_van`. Features we are dealing with:

1. Mileage
2. Vehicle condition
3. Year of registration
4. Body type
5. Fuel type

We are working on to predict the price of a car, so price is our target and all above mentioned features are predictors.

In [410]:

```
test=advert.drop(columns=['public_reference', 'reg_code', 'standard_colour', 'standard_make', 'standard_model', 'crossover_car_and_van'], axis=1)
```

In [411]:

```
sample_advert.isna().sum()
```

Out[411]:

```
mileage           127
vehicle_condition    0
year_of_registration 33311
price              0
body_type          837
fuel_type          601
dtype: int64
```

Above code shows the sum of null values in our potential features, so now we'll deal with the missing values in those features. Data is important and we can't afford losing any of it. However we are not using `drop()` function neither to drop a column nor a row.

1. For Mileage, year of registration and price we are replacing the null values with the mean of the entire column using `fillna()` and `mean()` function.

In [412]:

```
sample_advert['mileage']=sample_advert['mileage'].fillna(sample_advert.mileage.mean())
```

In [413]:

```
sample_advert['year_of_registration']=sample_advert['year_of_registration'].fillna(sample_advert.year_of_registration.mean()).astype(int)
```

In [414]:

```
sample_advert['price']=sample_advert['price'].fillna(sample_advert.price.mean())
```

2. To deal with missing values in `body_type` and `fuel_type`, we have replaced the null values with the maximum number of observations in that columns. For eg. In case of `body_type`, maximum number of `body_type` is 'Hatchback', so we have replaced it with 'Hatchback'.

In [415]:

```
sample_advert['body_type'].value_counts()
```

Out[415]:

```
Hatchback      167315
SUV             115872
Saloon          36641
Estate         24692
Coupe           23258
Convertible     16038
MPV             16026
Pickup           620
Combi Van       214
Limousine       159
Minibus         149
Camper           77
Panel Van        61
Window Van       41
Chassis Cab       3
Car Derived Van   2
Name: body_type, dtype: int64
```

In [416]:

```
sample_advert['body_type']=sample_advert['body_type'].fillna("Hatchback")
```

In [417]:

```
sample_advert['fuel_type'].value_counts()
```

Out[417]:

```
Petrol          216929
Diesel          158120
Petrol Hybrid   13602
Petrol Plug-in Hybrid 6160
Electric        4783
Diesel Hybrid   1403
Bi Fuel         221
Diesel Plug-in Hybrid 185
Natural Gas      1
Name: fuel_type, dtype: int64
```

In [418]:

```
sample_advert['fuel_type']=sample_advert['fuel_type'].fillna("Petrol")
```

In [419]:

```
sample_advert.isna().sum()
```

Out[419]:

```
mileage          0
vehicle_condition 0
year_of_registration 0
price            0
body_type        0
fuel_type        0
dtype: int64
```

Now, as we can see all the missing values are removed.

2.2 Handling outliers

To deal with the outliers, we can use two techniques Z-score and Interquantile range (IQR).

IQR= Q3-Q1 It is basically the first quartile subtracted from the third quartile and measures the dispersion similar to standard deviation.

Here we are dealing with the outliers in year_of_registration.

In [420]:

```
sample_advert['year_of_registration'].describe()
```

Out[420]:

```
count    402005.000000
mean      2015.005691
std         7.625632
min        999.000000
25%       2014.000000
50%       2016.000000
75%       2018.000000
max       2020.000000
Name: year_of_registration, dtype: float64
```

In [421]:

```
upper_boundary=sample_advert['year_of_registration'].mean()+3*sample_advert['year_of_registration'].std()
lower_boundary=sample_advert['year_of_registration'].mean()-3*sample_advert['year_of_registration'].std()
print(lower_boundary),print(upper_boundary),print(sample_advert['year_of_registration'].mean())
```

```
1992.1287952316097
2037.8825877113886
2015.0056914714992
```

Out[421]:

(None, None, None)

In [422]:

```
IQR=sample_advert.year_of_registration.quantile(0.75)-sample_advert.year_of_registration.quantile(0.25)
```

In [423]:

```
lower_bridge=sample_advert['year_of_registration'].quantile(0.25)-(IQR*1.5)
upper_bridge=sample_advert['year_of_registration'].quantile(0.75)+(IQR*1.5)
print(lower_bridge),print(upper_bridge)
```

```
2008.0
2024.0
```

Out[423]:

(None, None)

After using the IQR technique, now we got the minimum and maximum year_of_registration. We can only work on those years.

In [424]:

```
sample_advert=sample_advert.query("year_of_registration >2007 and year_of_registration<2024")
```

In [426]:


```
sns.boxplot(sample_advert['year_of_registration'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[426]:

```
<AxesSubplot: xlabel='year_of_registration'>
```



The outliers in year_of_registration is removed as seen in the boxplot above.

In [427]:

```
sample_advert['price'].describe()
```

Out[427]:

```
count    3.786670e+05
mean     1.790488e+04
std       4.693150e+04
min       2.500000e+02
25%       8.000000e+03
50%      1.300000e+04
75%      2.069500e+04
max       9.999999e+06
Name: price, dtype: float64
```

In [428]:

```
upper_boundary=sample_advert['price'].mean()+3*sample_advert['price'].std()
lower_boundary=sample_advert['price'].mean()-3*sample_advert['price'].std()
print(lower_boundary),print(upper_boundary),print(sample_advert['price'].mean())
```

```
-122889.60707963898
158699.36707456858
17904.87999746479
```

Out[428]:

(None, None, None)

In [429]:

```
IQR=sample_advert.price.quantile(0.75)-sample_advert.price.quantile(0.25)
```

In [430]:

```
lower_bridge=sample_advert['price'].quantile(0.25)-(IQR*1.5)
upper_bridge=sample_advert['price'].quantile(0.75)+(IQR*1.5)
print(lower_bridge),print(upper_bridge)
```

```
-11042.5
39737.5
```

Out[430]:


(None, None)

In [432]:

```
sns.scatterplot(data=sample_advert['price'])
```

Out[432]:

<AxesSubplot:ylabel='price'>



2.3 Feature engineering

In this part, we will be working on converting the categorical data to numeric, to check the correlation between all the features (predictors).

1. Firstly we will scale the data, where it gives structure to data and standardize the range of features of an input data set.

In [433]:

```
from sklearn.preprocessing import StandardScaler
```

In [434]:

```
Scaler=StandardScaler()
```


In [436]:

```
Scaler.fit_transform(sample_advert[['mileage', 'year_of_registration']])
```

Out[436]:

```
array([[ -1.08346843,  -0.25340458],
       [  2.31985881,  -1.6303757 ],
       [ -0.83819491,   0.43508098],
       ...,
       [  0.57558678,  -1.28613292],
       [ -0.76115388,  -0.25340458],
       [ -0.64323392,  -0.59764736]])
```

In [463]:

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, OrdinalEncoder
oe=OrdinalEncoder()
```

We are using Ordinal Encoding for vehicle_condition, body_type and fuel_type

In [464]:

```
sample_advert.vehicle_condition=oe.fit_transform(sample_advert[["vehicle_condition"]]).astype(int)
```

In [465]:

```
sample_advert.body_type=oe.fit_transform(sample_advert[["body_type"]]).astype(int)
```

In [466]:

```
sample_advert.fuel_type=oe.fit_transform(sample_advert[["fuel_type"]]).astype(int)
```

After scaling the data and converting categorical values to numerical we have the dataset as below:

In [441]:

```
sample_advert.head()
```

Out[441]:

	mileage	vehicle_condition	year_of_registration	price	body_type	fuel_type
0	0.0	0	2015	73970	13	7
1	108230.0	1	2011	7000	14	1
2	7800.0	1	2017	14000	13	5
3	45000.0	1	2016	7995	7	1
4	64000.0	1	2015	26995	13	1

In [442]:

```
sample_advert.corr()['price']
```

Out[442]:

```
mileage          -0.152224
vehicle_condition -0.095481
year_of_registration  0.118914
price             1.000000
body_type         0.043248
fuel_type         0.018249
Name: price, dtype: float64
```

Now again, checking the correlation between all the features (predictors) and price (target). Year_of_registration, body_type and fuel type shows positive and correlation, while mileage shows the least.

Section-3 Analysis


1.1 Quantitative - Quantitative analysis

In [469]:

```
sns.lmplot(
    data=sample_advert, x='year_of_registration', y='price', col='body_type', col_wrap=4,
    scatter_kws=dict(alpha=0.1), height=4, lowess=True).set(ylim=(0,300000))
```

Out[469]:

<seaborn.axisgrid.FacetGrid at 0x216e4f8d670>



Observations:


1. Respective to body_type the price of a car increases/ decreases along with the year_of_registration.
2. Body type 8 has a huge spike in the price in the year between 2018-19.
3. On the other hand, body type 3 has a fall in the price has the year increased.

In [447]:

```
sns.barplot(data=sample_advert, x='year_of_registration', y='price')
```

Out[447]:

<AxesSubplot: xlabel='year_of_registration', ylabel='price'>



Observations:

1. Price increases as year of registration increases.
2. Spike in price in year 2015.


3.2 Quantitative-Categorical analysis

In [352]:

```
price_bodytype=sns.barplot(data=advert, x='body_type',y='price').set(ylim=(0,100000))
price_bodytype
plt.xticks(rotation=45)
```

Out[352]:

```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15]),
 [Text(0, 0, 'SUV'),
  Text(1, 0, 'Saloon'),
  Text(2, 0, 'Hatchback'),
  Text(3, 0, 'Convertible'),
  Text(4, 0, 'Limousine'),
  Text(5, 0, 'Estate'),
  Text(6, 0, 'MPV'),
  Text(7, 0, 'Coupe'),
  Text(8, 0, 'Pickup'),
  Text(9, 0, 'Combi Van'),
  Text(10, 0, 'Panel Van'),
  Text(11, 0, 'Minibus'),
  Text(12, 0, 'Window Van'),
  Text(13, 0, 'Camper'),
  Text(14, 0, 'Car Derived Van'),
  Text(15, 0, 'Chassis Cab')])
```




Observations:

1. Limousine is the most expensive body type, camper takes the second spot.
2. Car Derived Van is the cheapest car.

In [353]:

```
price_fueltype=sns.barplot(data=advert, x='fuel_type',y='price',color='purple')  
price_fueltype  
plt.xticks(rotation=45)
```

```
Text(8, 0, 'Natural Gas']])
```



Observations:


1. Diesel Hybrid is the most expensive, petrol plug-in hybrid and diesel plug-in hybrid are the second most expensive.
2. Natural Gas is the cheapest fuel_type.

In [354]:

```
sns.barplot(x = 'vehicle_condition', y = 'price', data = advert)
```

Out[354]:

```
<AxesSubplot:xlabel='vehicle_condition', ylabel='price'>
```



Observations:


1. New cars are more expensive than used cars.

In [309]:

```
advert['fuel_type'].value_counts(normalize=True).plot.bar()
```

Out[309]:

<AxesSubplot:>




In [310]:

```
advert['body_type'].value_counts(normalize=True).plot.bar()
```

Out[310]:

<AxesSubplot:>



3.3 Categorical-Categorical analysis

In [474]:


```
rnd = advert.groupby(['body_type', 'fuel_type'])['price'].median().unstack()
rnd.head(10)
```

Out[474]:

	fuel_type	Bi Fuel	Diesel	Diesel Hybrid	Diesel Plug-in Hybrid	Electric	Natural Gas	Petrol	Petrol Hybrid	Petrol Plug-in Hybrid
body_type										
	Camper	9497.5	39497.5	NaN	NaN	NaN	NaN	13750.0	NaN	NaN
	Car Derived Van	NaN	7245.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
	Chassis Cab	NaN	8750.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
	Combi Van	NaN	14995.0	NaN	NaN	30933.5	3795.0	15995.0	NaN	NaN
	Convertible	57000.0	14295.0	NaN	NaN	18944.0	NaN	14990.0	1099950.0	69147.0
	Coupe	41995.0	14995.0	47565.0	NaN	14495.0	NaN	22995.0	27295.0	48500.0
	Estate	3790.0	11950.0	31221.0	19500.0	25216.5	NaN	15990.0	14799.5	27997.0
	Hatchback	11519.5	9495.0	8297.0	NaN	23990.0	NaN	8495.0	13295.0	21495.0
	Limousine	NaN	8497.0	NaN	NaN	NaN	NaN	59995.0	NaN	NaN
	MPV	5750.0	9995.0	10550.0	NaN	32475.0	NaN	9000.0	19397.5	20995.0

In [475]:

```
plt.subplots(figsize=(8,6))
sns.heatmap(
    data=rnd, cmap='Greens'
);
```



Observations:

1. Convertible with Petrol hybrid is most expensive.

In [313]:

```
avg_price=advert.groupby(['body_type', 'vehicle_condition'])['price'].mean()
```


In [314]:

```
avg_price
```

Coupe	USED	26387.292081
	NEW	55624.437666
	USED	35041.950898
Estate	NEW	37462.720141
	USED	14866.036989
Hatchback	NEW	20993.277737
	USED	9709.684775
Limousine	USED	58953.911950
MPV	NEW	29510.039039
	USED	10828.357803
Minibus	NEW	39935.000000
	USED	20179.931507
Panel Van	USED	18028.868852
Pickup	NEW	30135.935484
	USED	19080.241935
SUV	NEW	38125.806350
	USED	20118.343669
Saloon	NEW	43461.319194
	USED	18014.629809
Window Van	USED	17240.853659

In [316]:

```
(advert
.groupby(['body_type', 'vehicle_condition'])
['price']
.mean()
.unstack()
.plot.bar()
);
```



Observations:

1. We do not have New cars in Camper, Car derived van, Chassis cab, Panel van, Limousine and window van.
2. Limousine (Used) the most expensive car.
3. Coupe (New) the most expensive car.

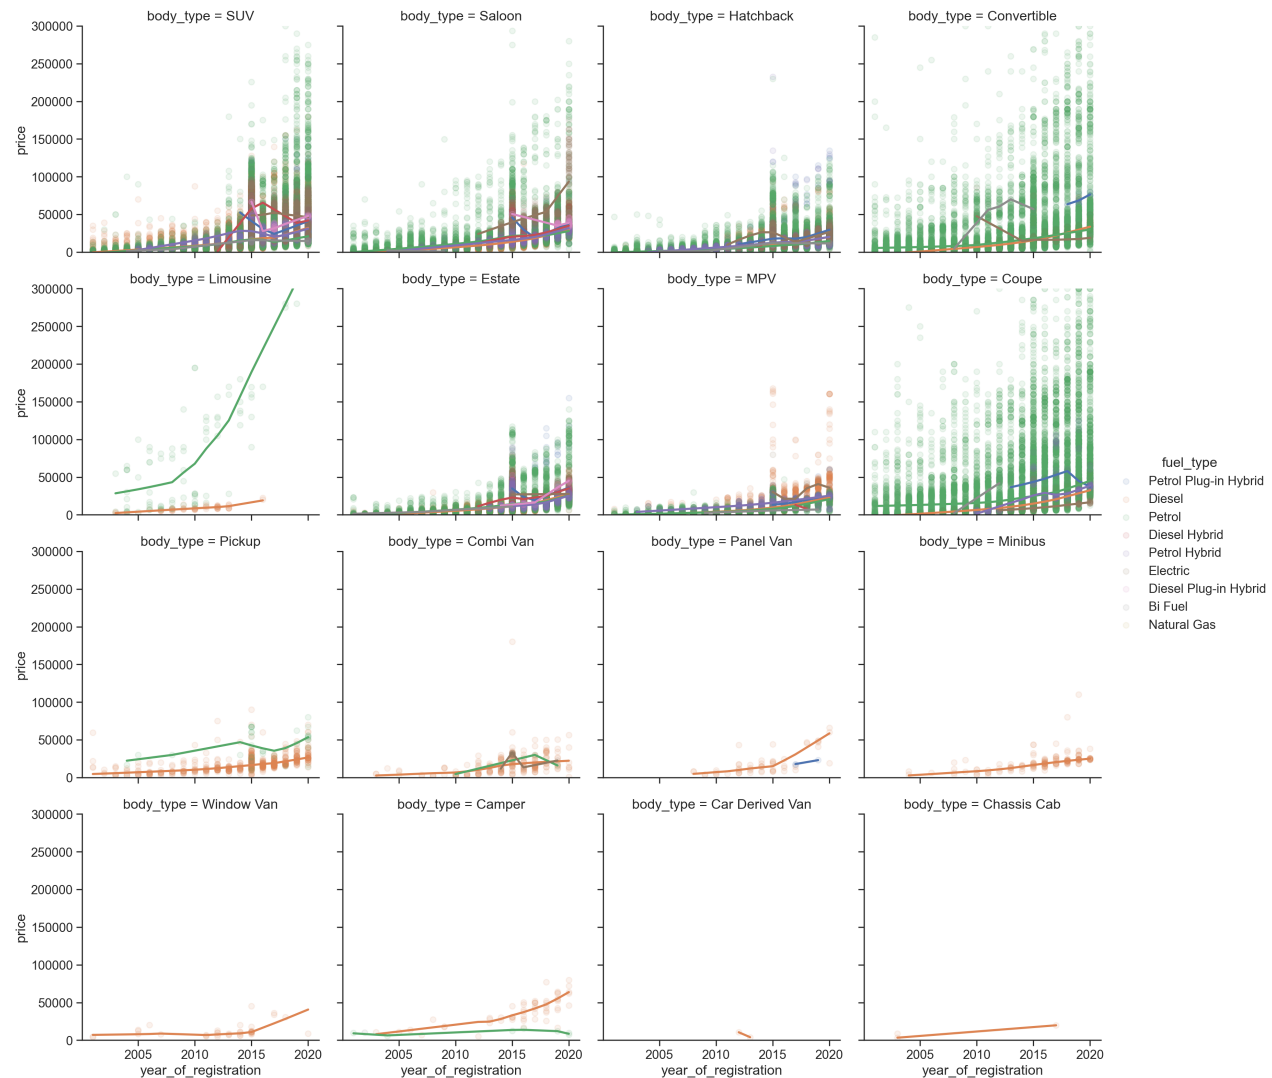
In [318]:

```
sns.lmplot(
    data=advert, x='year_of_registration', y='price', hue='fuel_type', col='body_type', col_wrap=4,
    scatter_kws=dict(alpha=0.1), height=4, lowess=True).set(ylim=(0,300000))
```

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\nonparametric\smoothers_lowess.py:227: RuntimeWarning: invalid value encountered in true_divide
 res, _ = _lowess(y, x, x, np.ones_like(x),

Out[318]:

<seaborn.axisgrid.FacetGrid at 0x2164cb80730>



In []:

In []: